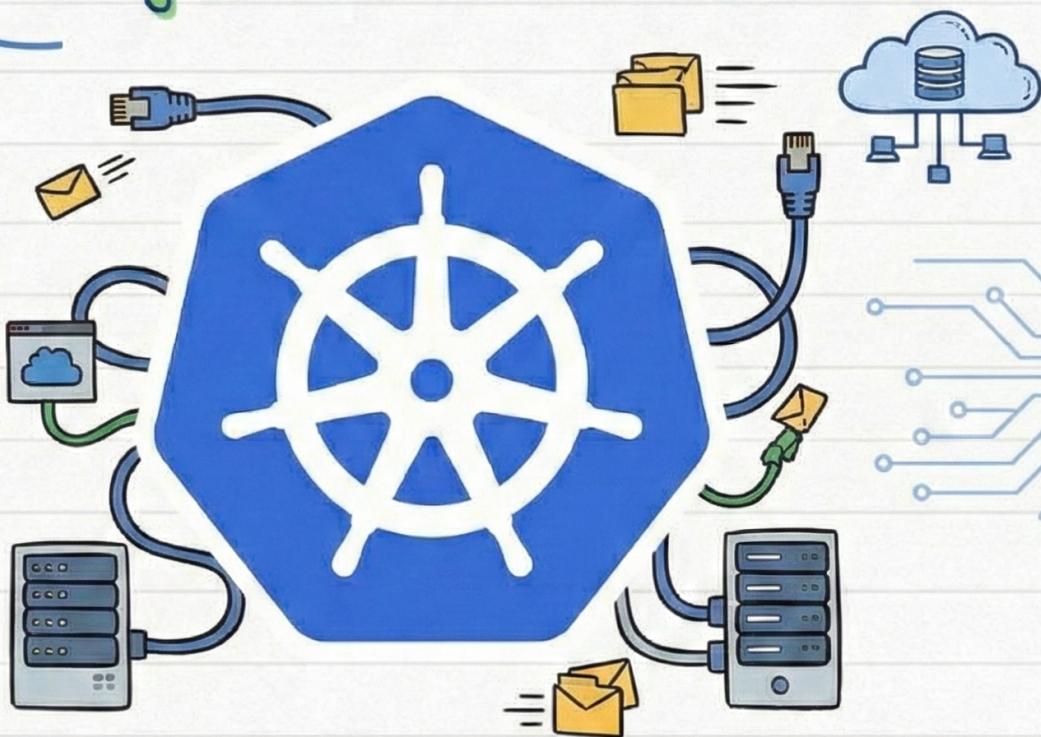


Kubernetes Simplified – Part 2



Handwritten By:
Vaibhav Matkar

Service → stable network endpoint used for communication.

service components

① **clusterIP** – default servicetype → make service reachability within cluster only. i.e exposes service to internal IP's [communication within cluster]

```
apiVersion: v1
kind: Service
metadata:
  labels:
    k8s-app: kube-dns
    kubernetes.io/cluster-service: "true"
    kubernetes.io/name: CoreDNS
  name: kube-dns
  namespace: kube-system
spec:
  clusterIP: 10.96.0.10
  ports:
  - name: dns
    port: 53
    protocol: UDP
    targetPort: 53
  - name: dns-tcp
    port: 53
    protocol: TCP
    targetPort: 53
  selector:
    k8s-app: kube-dns
  type: ClusterIP
```

② **Nodeport** - Expose service outside the cluster by opening port on every node. [communication outside cluster]

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  type: NodePort
  selector:
    app.kubernetes.io/name: MyApp
  ports:
    - port: 80
      # By default and for convenience, the `targetPort` is set to
      # the same value as the `port` field.
      targetPort: 80
      # Optional field
      # By default and for convenience, the Kubernetes control plane
      # will allocate a port from a range (default: 30000–32767)
      nodePort: 30007
```

③ **Externalname** → just maps service.name to external DNS

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
  namespace: prod
spec:
  type: ExternalName
  externalName: my.database.example.com
```

④ **Loadbalancer** → created by cloud provider
exposes service to internet with public ip / DNS

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app.kubernetes.io/name: MyApp
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
  clusterIP: 10.0.171.239
  type: LoadBalancer
status:
  loadBalancer:
    ingress:
      - ip: 192.0.2.127
```

Network communication in kubernetes

- ① container to container → via local host as
belongs to same network namespace
- ② pod to pod within same Node → Each pod has
its own ip so communicate via IPs using CNI
- ③ pod to pod across different Nodes (same cluster)
→ CNI plugin setup routing/overlay network
across nodes to communicate

Kubernetes Network Security & Access control

- ① Network policies → rules that restricts traffic inside cluster
- ② RBAC → granular permissions for users & services
- ③ CNI → plugin based interface allows 3rd party network providers to give additional network security

Ingress in kubernetes

↳ App layer (7th)

↓
mgmt of incoming traffic to cluster
using adv routing & rules for HTTP/HTTPS

Functions

① Load Balancing

↓
distribute traffic to different instances of service.

② Routing

↓
distribute traffic to multiple services correctly using routing rules

③ Access control

↓
prevents un-autho-access to any service based on ip.

④ No need of TLS certification for every service instead use in ingress

Encr → Decry

Working of Nodeport service



Allows to expose container/application to outside world

i> create Nodeport service in k8 manifest file

mention ↗ Target port → container listening
↗ port to be exposed to outside world
ie. Nodeport

↓
manually/auto [30,000 - 32767]

ii> Kubernetes maps Target port to Nodeport

iii> To access application outside Kubernetes cluster → `http://node-ip/nodeport`

kube proxy

- `kube proxy` → Agent installed on every node ↗ as daemonset ↗ Linux process

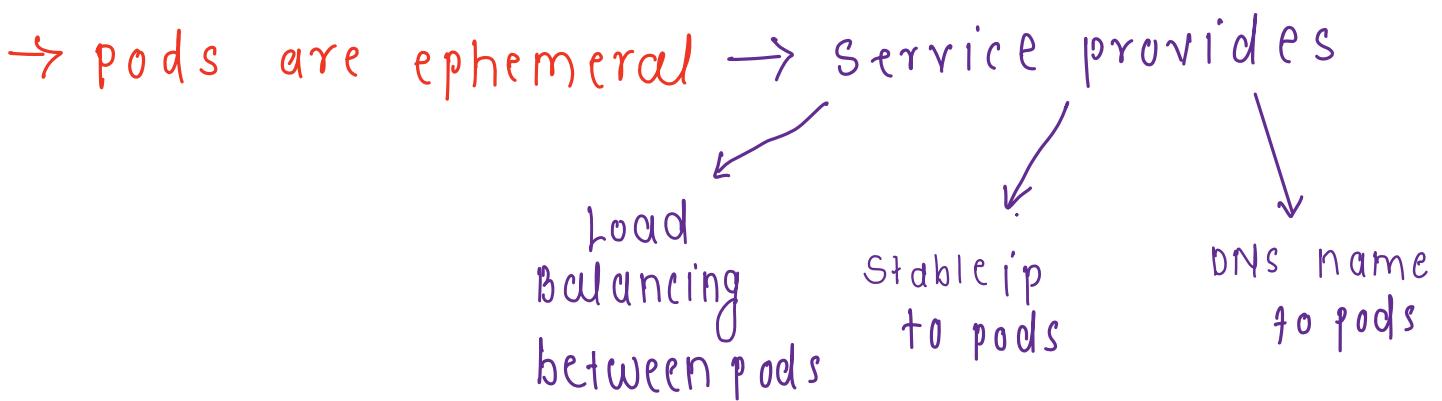
middleware between service & pod

• function → monitor changes in service and

endpoints [pod ip's] → create network rules



Decides which request will go to which pod



Working of kube proxy

- ① Any new service / endpoints added → API server inform this to kube proxy
- ② kube proxy create NAT rule for this changes.
i.e. map service ip to pod ip → when request come to service it is redirected to pod

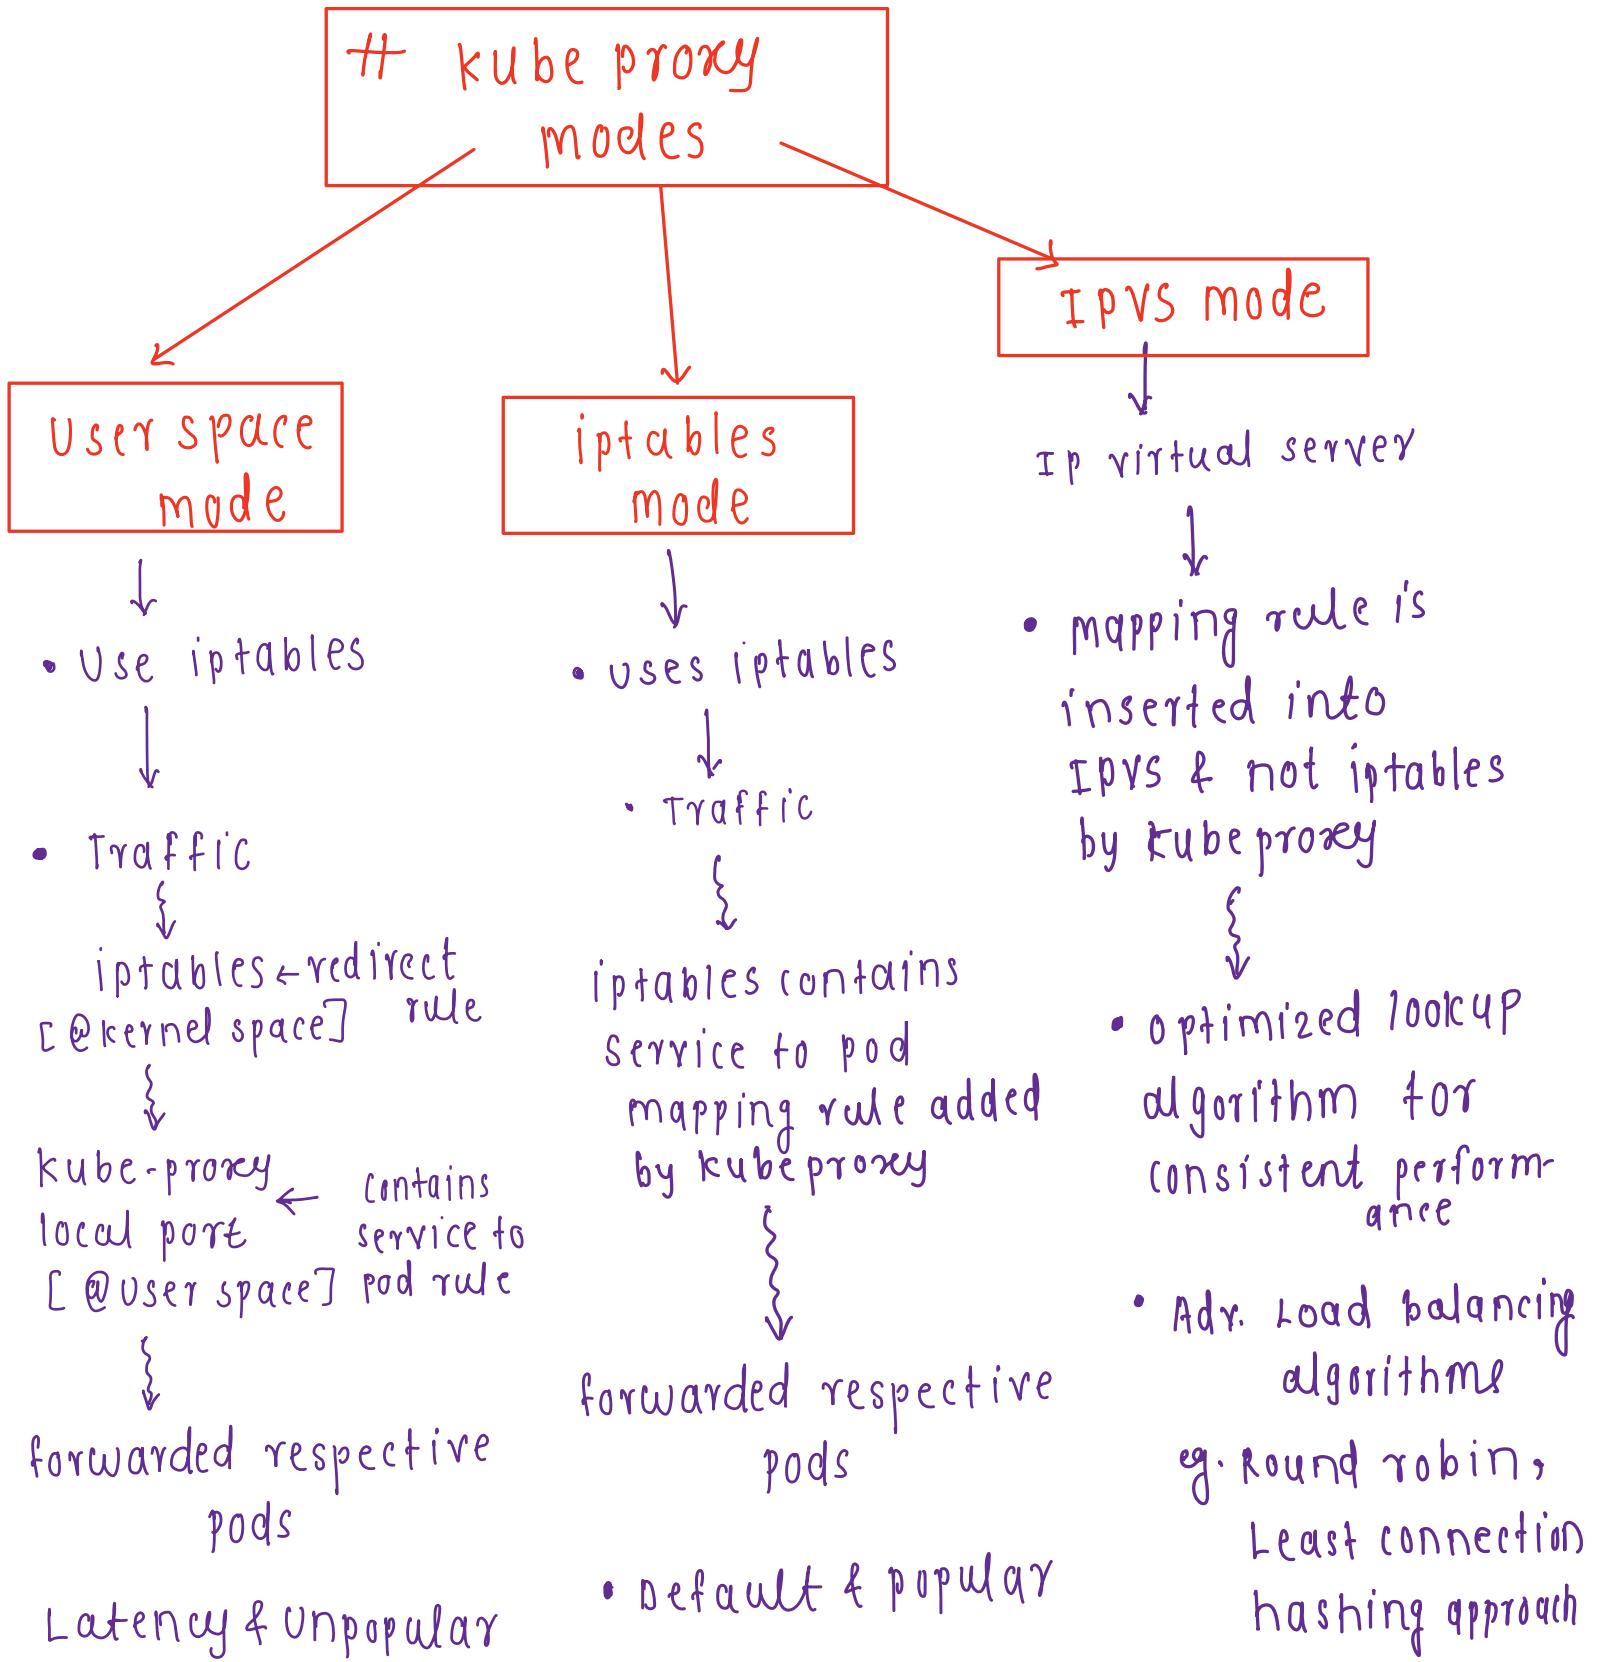
NAT [Network Address translation]

↓
Rule that forwards ip address/port (1st)
to ip address/port (2nd)

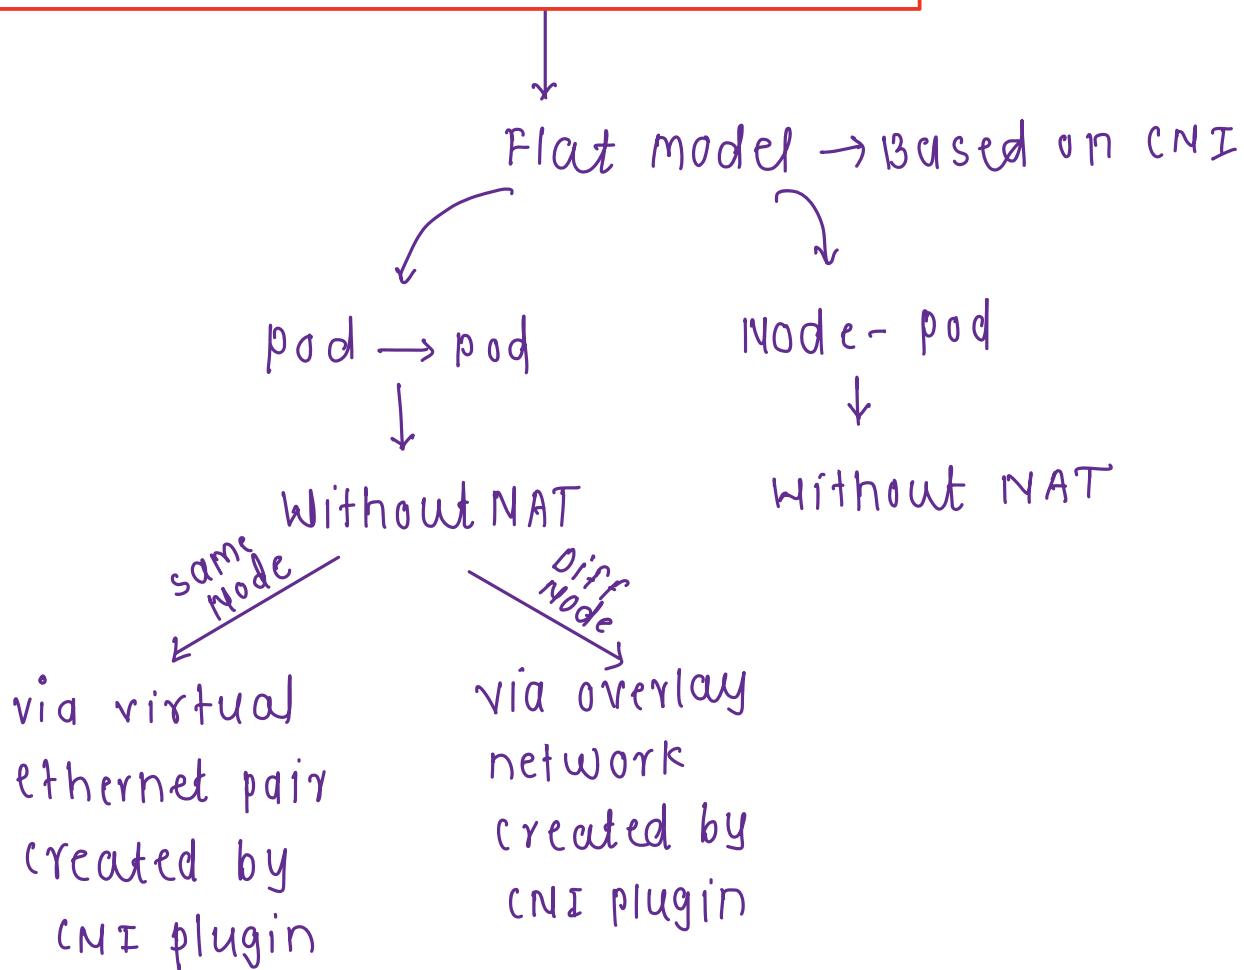


Internal private IP → public IP

redirect external request to appropriate destination

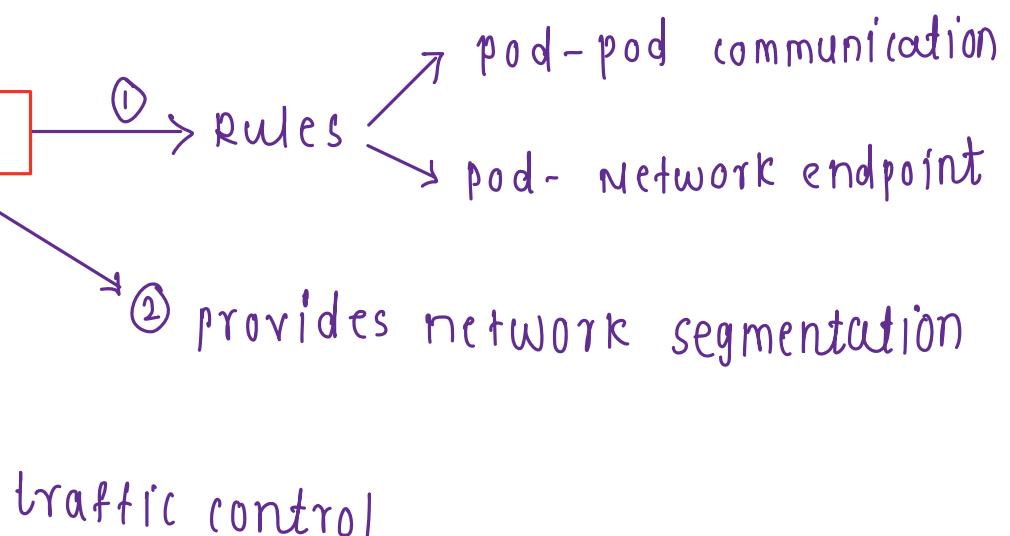


Kubernetes networking model



Network policy

Restrict access
to sensitive
services



kubernetes dns

↳ [internal DNS] for service discovery

- pod → DNS A record form → pod-ip-addr.namespace.pod.cluster.local
- service → DNS A record → service-name.namespace.svc.cluster.local

challenges of implementing k8s network policies.

- ① Larger clusters => more policy complexity
- ② Policy conflicts => more than 1 policy for same pod
- ③ performance ↓ => cluster large consumes more resources during policy evaluation leaving few resources.

troubleshooting network connectivity issues in k8s cluster

- ① verify CNI configuration → directory → /etc/cni/net.d/
- ② check pod label & selector → incorrect network policy applied to right pod
- ③ look for policy conflicts → nx policy → single pod
- ④ check CNI log files → vary → /var/log/cni plug/cni.log

Egress Network policy

→ pods are allowed to send traffic outside

Ingress Network policy.

→ pods are allowed for incoming traffic

- podselector → select pod matching label

podselectors:

matchLabels:

app: demo

- namespaceSelector → select all pod matching label

namespaceSelector:

matchLabels:

app: demo

- ipBlock → allow traffic from specific CIDR range

ipBlock:

cidr: <ip>/24

- from: [] → rules used for ingress policy
↳ default

- to: [] → rules used for egress policy
↳ default

#Container Network Interface [CNI]

↳ plugin system that decides networking in kubernetes.

- Flannel → overlay network provider
- calico → provides secure network connectivity

Working of flannel

- Flannel creates flat network
- Assign subnet to each node
 - ↓
 - Each pod gets ip from subnet
- Flanneld agent maintains routing table mapping subnets to node IP's.

- communication
 - Pod-pod (same Node) → via linux bridge [No encapsulation]
 - Pod-pod (Across Node) → via VXLAN encapsulation by flanneld of one node & decapsulation by flanneld of other Node.

• **+ calico CNI**

↳ open source networking & network security solution

• **calico Architecture**

- ① **Birds** → BGP routing daemon that distribute routes from felix to BGP peers [nodes]
- ② **Felix** → manages routes/ACL's & other configurations on each host [Node]
- ③ **confd** → Tool that monitors changes in calico configuration in datastore.
- ④ **IPAM** [IP Address mgmt] → Assign ip's to pods using IP Pool resources
- ⑤ **Typha** → cache between datastore & felix instance
- ⑥ **Datastore** → stores calico related configurations

• **Working of calico**

- New pod $\xrightarrow{\text{calls}}$ calico → Assign unique ip to pod & add pod to linux interface
- Each Node has CIDR range → calico adds route to node's kernel routing table.

- calico allows for network policies for pod communication.

- calico uses pure layer 3 routing → packet transfers from node to node without the encapsulation.



USES BGP [Border Gateway protocol]



Nodes shares routes with each other, so each node knows which node has which pod.

calico Networking Modes

- ① IPIP mode → package encapsulation for inter-node communication.
- ② Direct mode → direct routing without encapsulation
- ③ VXLAN mode → package encapsulation for inter-node communication via vxlan interface
- ④ pod packet Encryption mode → uses WireGuard tunnels & encrypt & transmit pod traffic betwn nodes.

#service discovery in k8's

- service registry → central repository where each service registers network locations/capabilities/metadata
- service discovery mechanism → Allows services to dynamically discover & locate other services based on service registry.

• Working

- ① Service registration → unique DNS Name & IP address is assigned to pod
- ② DNS resolution → DNS resolution by other services.
- ③ Load balancing → during multiple requests for service, K8 automatically load balances traffic to available pod belong to service.
- ④ Service updates → if pods associated with services changes due to scalability → K8 updates service's DNS record for set of available endpoints.

NFS → protocol which allows to mount storage device as local drive on container

- Benefits
 - use existing storage volume
 - persistence
 - used to share data between containers
 - simultaneous mounting → NFS can be mounted on multiple nodes → R/W at same time from multiple nodes.

MTLS in kubernetes

↳ Mutual Transport Layer Security.



- zero trust mechanism → Both client & server verify each other using certificates and communicate

