

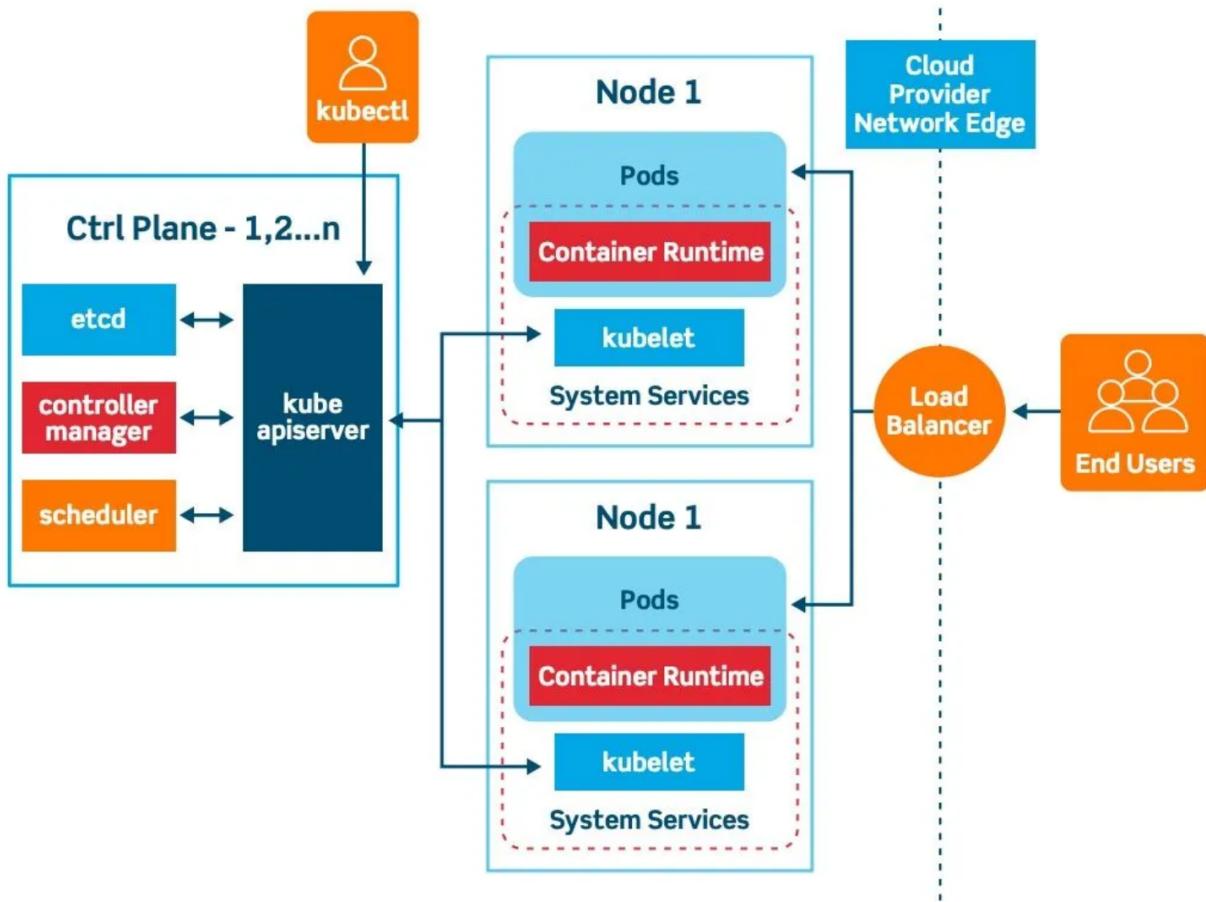
Kubernetes Simplified - Part 1



Handwritten By:
Vaibhav Matkar

kubernetes → container orchestration tool

kubernetes Architecture



- kubernetes cluster → set of nodes
 - Master Node - contains control plane → mgmt of cluster
 - Worker Node - Machines where containers actually run
- control plane → s/w component → mgmt of state of cluster

Components of master node

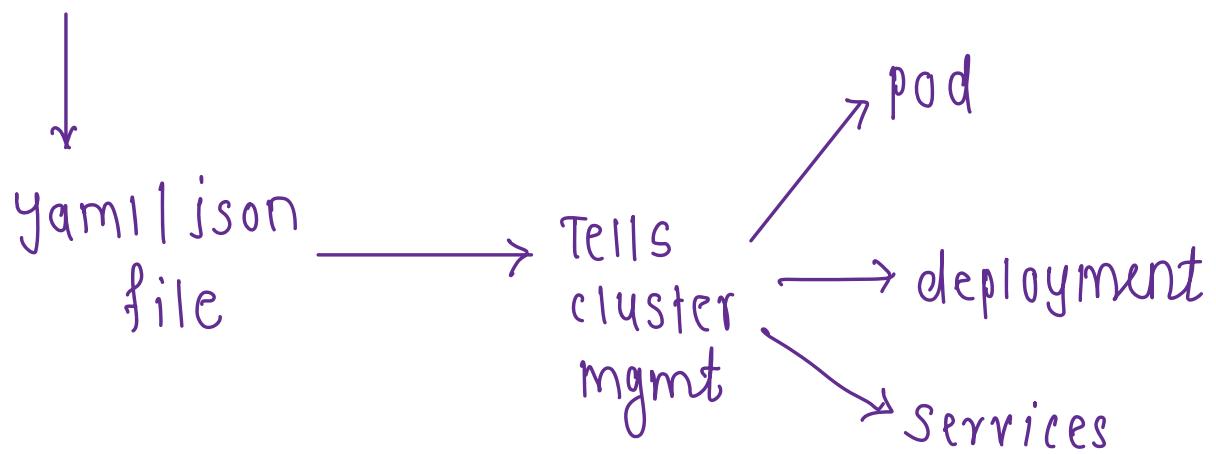
- ① API server → Act as frontend to cluster
 - ↳ Nerves of kubernetes → responsible for communication
- ② ETCD → stores cluster data in key-value store
 - ↳ Brain of the kubernetes.
- ③ control Manager → Manages different controllers
 - ① Node controller → Notice + respond if down
 - ② Replication controller → ensures specified number of pod replica's are running or not
 - ③ Endpoints controller → connects service & pod using labels.
 - ④ serviceAccount Token controller → Watches serviceAccounts → creates secret containing token → Mount token into pods running under that service account.
- ④ kube-scheduler → Assign each pod to a suitable node

Components of node

- ① kubelet → makes sure containers are running in pod with help of pods specs
 - Agent
 - It does not manage containers not created by k8's
 - ② kube-proxy → maintains network rules on nodes
 - ie. Load balancing between pods
 - ③ container-runtime → software responsible for running containers
 - ④ pods → smallest deployable unit in kubernetes.

`# kubectl` → cmdline tool to interact with k8 cluster.

manifest file in kubernetes



↑ components of manifest file

- ① apiVersion – Usable kubernetes API version
- ② kind – type of kubernetes resource/object
- ③ metadata - identifying data about resource
 - i) name → unique name given to resource
 - ii) namespace → opt
 - iii) Labels → key-value pairs
 - iv) Annotations → non-identifying meta-data
- ④ Spec – gives kubernetes desired state of the resource/ object
- ⑤ Status – gives current state of the resource

```
1  apiVersion: v1
2  kind: Pod
3  √ metadata:
4    |   name: web-application
5  √ spec:
6    √   containers:
7      √     - name: web-app
8        |   image: web-app-devansh
9      √     env:
10     √       - name: APP-COLOR
11     √         valueFrom:
12       √           configMapKeyRef:
13         |   name: app-config
14           |   key: APP-COLOR
```

Node Affinity → decides on which node pod can run

pod Affinity → decides pod-pod relationship i.e which pod should run with which another pod

Taint → Reserve node for special workload otherwise repel pods

Toleration → property of pods that allows them to run on node if their capacity matches taint.

Ingress controller



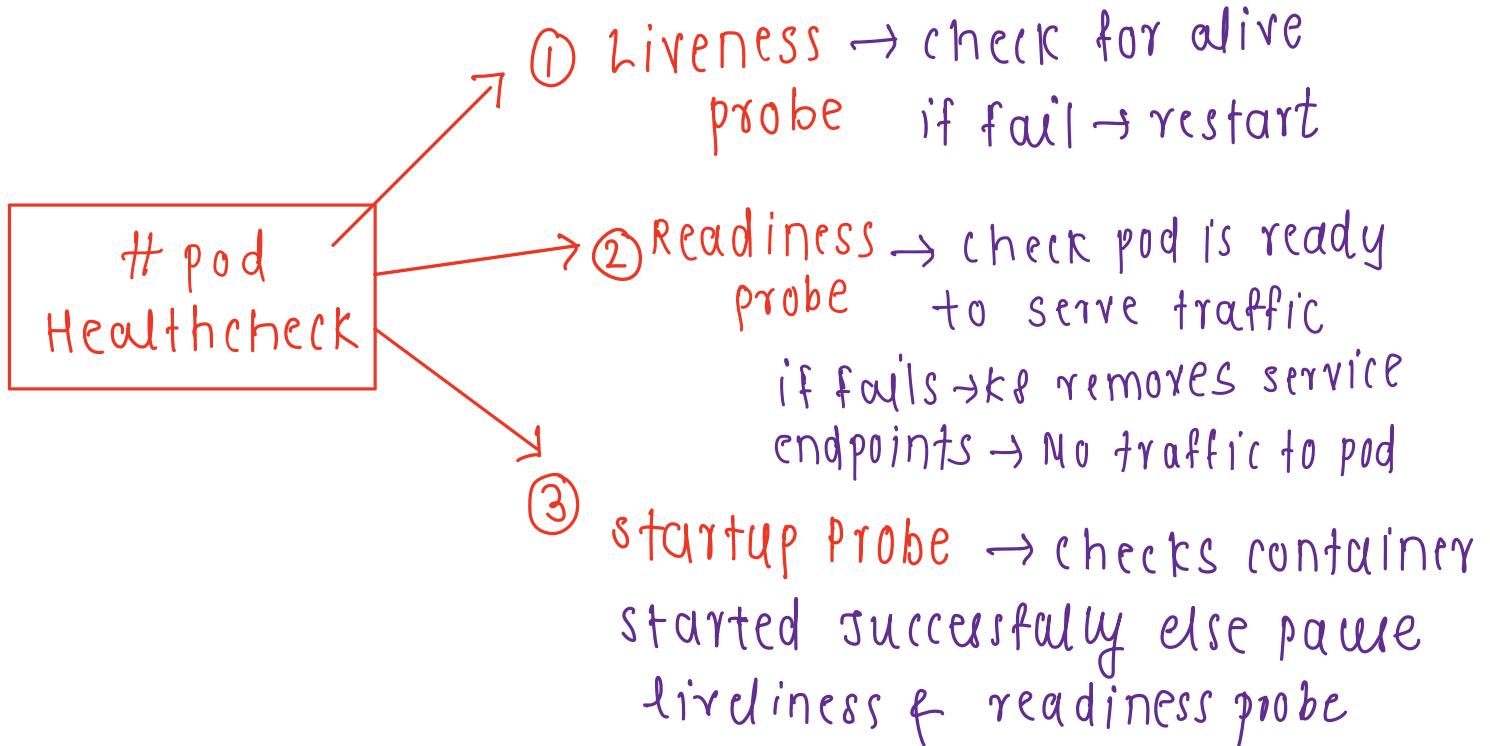
slow component that decides rules for how external traffic reach services inside cluster

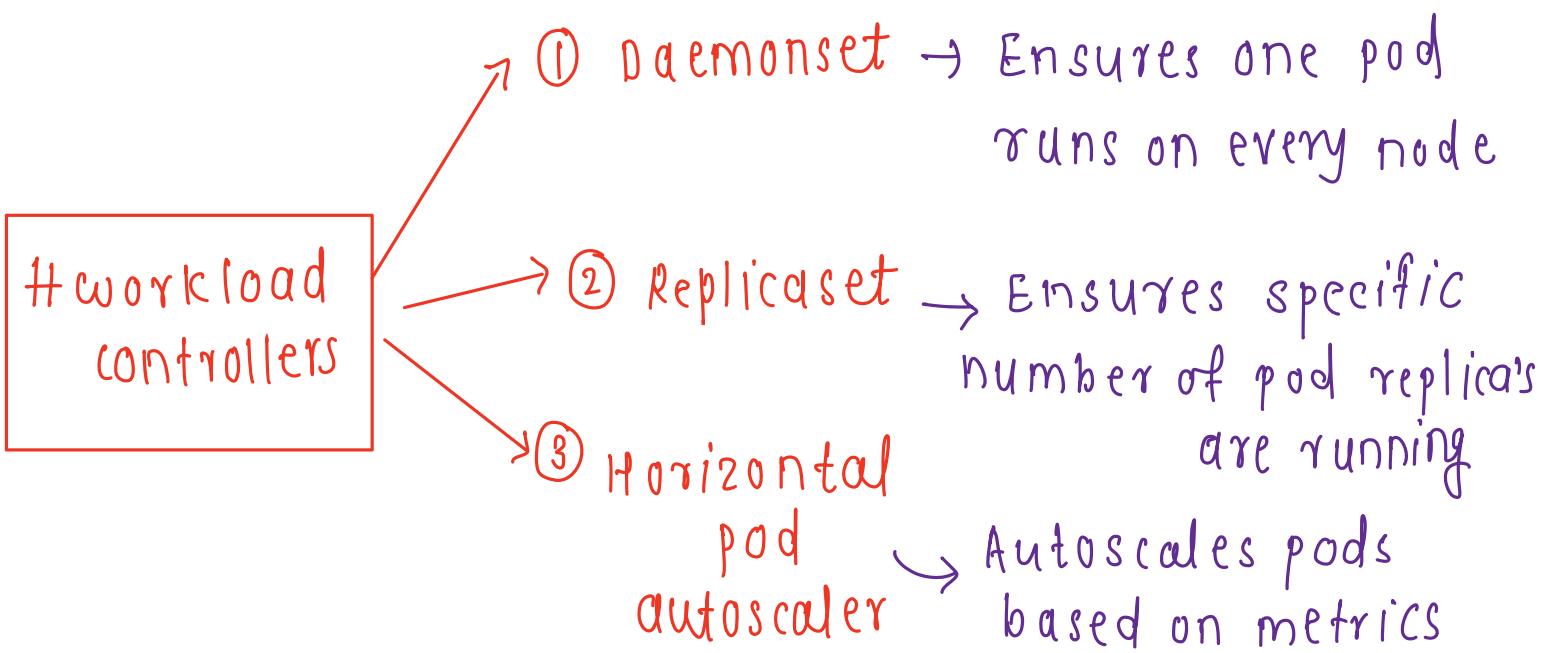
Namespace → Allows to separate resources without conflict eg. dev/ QA/ prod

context → parameters / shortcut that tells kubectl which cluster, user & namespace to use

Lifecycle of a pod

- ① pending → pods created | containers not running)
scheduler finding node to place pods
- ② Running → pod scheduled on node | containers
are created & running
- ③ succeeded → All containers exited successfully
| pod does not restart again
- ④ failed → Atleast one container terminated
with failure | pods won't restart auto.
- ⑤ Unknown → kubernetes can't communicate with
node (Network error)

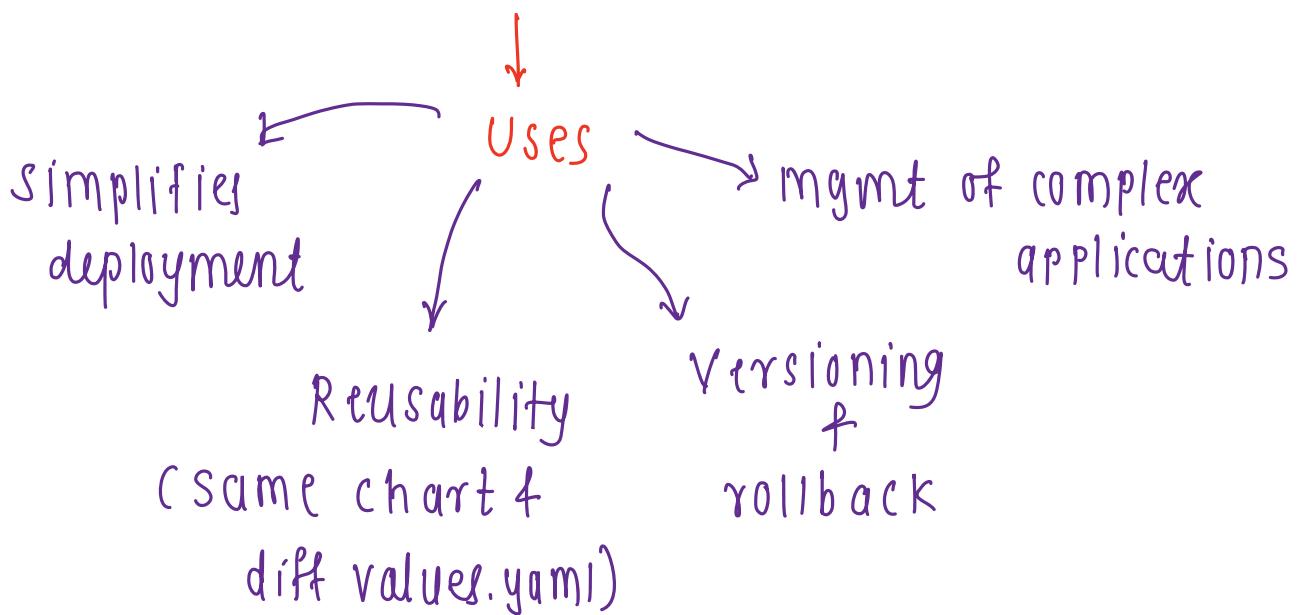




Helm → package manager for kubernetes

#Helm repository → collection of charts for kubernetes packages.

#Helmchart → contains all necessary resources & configurations to deploy application/service on k8s



kubernetes object → persistent entities within kubernetes that defines desired state of cluster.
eg. pods / deployment / replicaset / statefulset.

Deployment in kubernetes

- ↳ Gives declarative updates to pods & replicasets
- ↳ common way to run applications

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx
spec:
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
  spec:
    containers:
      - name: nginx
        image: nginx
        resources:
          limits:
            memory: "128Mi"
            cpu: "500m"
        ports:
          - containerPort: 80
```

Replicaset in kubernetes

- ↳ Ensures specified number of pods are running at all times.

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: nginx-replicaset
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx-rs-pod
  matchExpressions:
  - key: env
    operator: In
    values:
    - dev
  template:
    metadata:
      labels:
        app: nginx-rs-pod
        env: dev
  spec:
    containers:
    - name: nginx
      image: nginx
      ports:
      - containerPort: 80
```

StatefulSet in Kubernetes → for stateful applications

having stable network identities & persistent storage.

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: web
spec:
  serviceName: "nginx"
  replicas: 4
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: registry.k8s.io/nginx-slim:0.8
          ports:
            - containerPort: 80
              name: web
      volumeMounts:
        - name: www
          mountPath: /usr/share/nginx/html
volumeClaimTemplates:
  - metadata:
      name: www
    spec:
      accessModes: [ "ReadWriteOnce" ]
      resources:
        requests:
          storage: 1Gi
```

Daemonset in Kubernetes

- ↳ Ensures one pod runs on every node.
- ↳ commonly used for logging/monitoring/ network agent

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: example-daemonset
spec:
  selector:
    matchLabels:
      app.kubernetes.io/name: example
  template:
    metadata:
      labels:
        app.kubernetes.io/name: example
    spec:
      containers:
        - name: pause
          image: registry.k8s.io/pause
      initContainers:
        - name: log-machine-id
          image: busybox:1.37
          command: ['sh', '-c', 'cat /etc/machine-id > /var/log/machine-id']
      volumeMounts:
        - name: machine-id
          mountPath: /etc/machine-id
          readOnly: true
        - name: log-dir
          mountPath: /var/log
      volumes:
        - name: machine-id
          hostPath:
            path: /etc/machine-id
            type: File
        - name: log-dir
          hostPath:
            path: /var/log
```

job in kubernetes

↳ runs the pod till completion / for batch processing

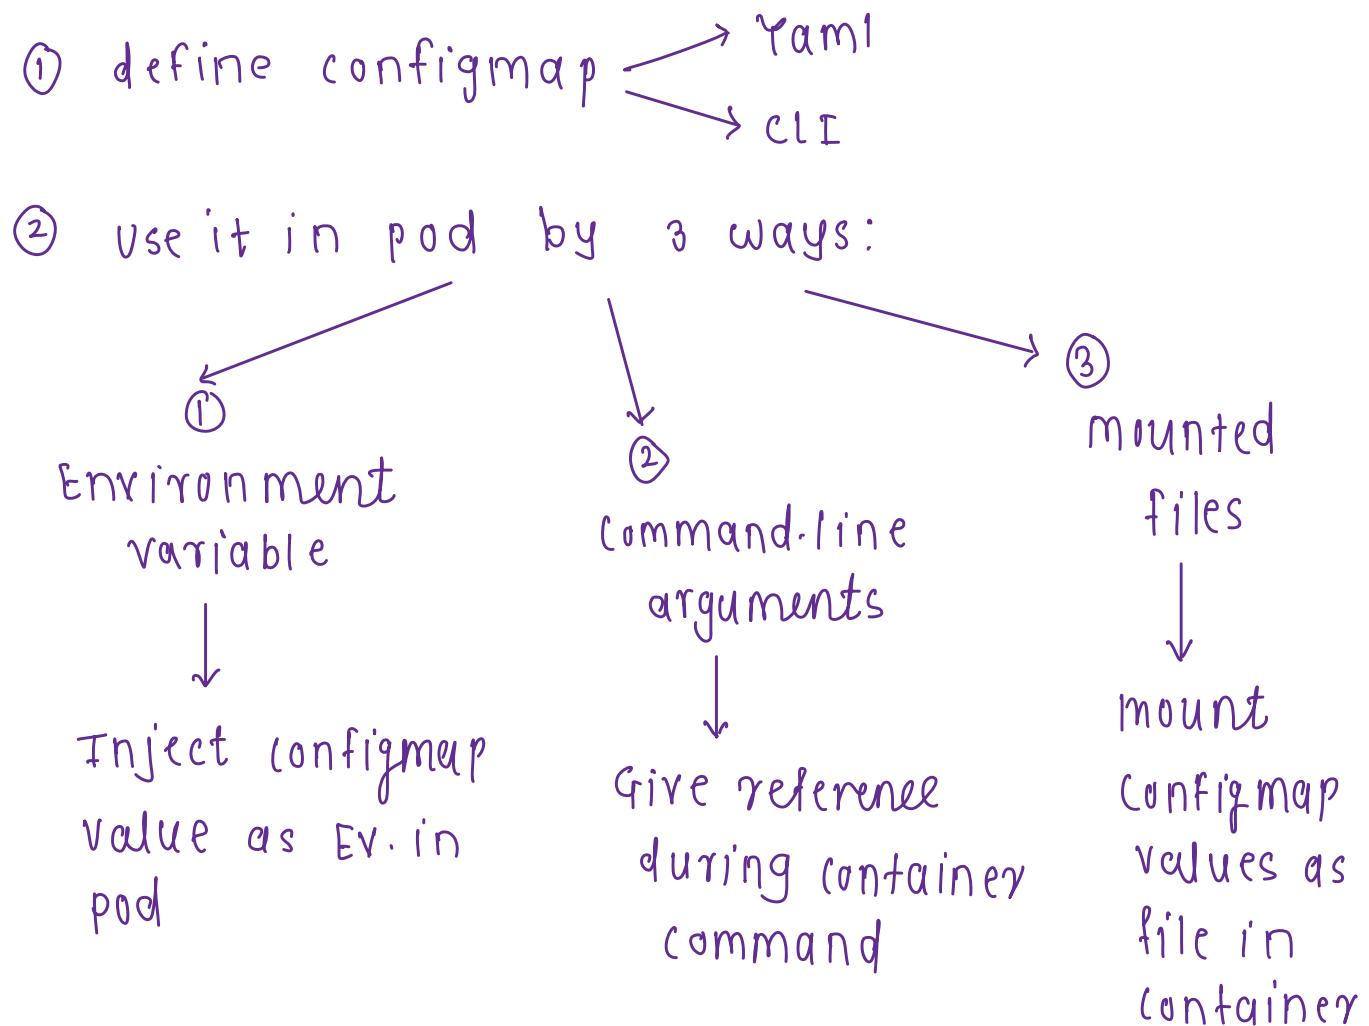
```
apiVersion: batch/v1
kind: Job
metadata:
  name: pi
spec:
  template:
    spec:
      containers:
        - name: pi
          image: perl:5.34.0
          command: ["perl", "-Mbignum=bpi", "-wle", "print bpi(2000)"]
      restartPolicy: Never
  backoffLimit: 4
```

cronjob in kubernetes

↳ runs job on schedule . eg . for backup script

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: hello
spec:
  schedule: "* * * * *"
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: hello
              image: busybox:1.28
              imagePullPolicy: IfNotPresent
              command:
                - /bin/sh
                - -c
                - date; echo Hello from the Kubernetes cluster
      restartPolicy: OnFailure
```

configmap working



Deploy pod in kubernetes

① using YAML manifest file

② using kubectl command directly

scaling a deployment

① manual scaling (kubectl scale command)

② declarative scaling

edit replicas in deployment.yaml

③ Autoscaling (HPA - based on cpu/memory)

Accounts in kubernetes

① User account → for users → To gain access to cluster
② service account → gives identity to pod to talk to k8 API



Namespace has default service account

creating a own service account:

① create a service account

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: cicd-deployer
  namespace: deployment
---
```

② create Role containing permissions.

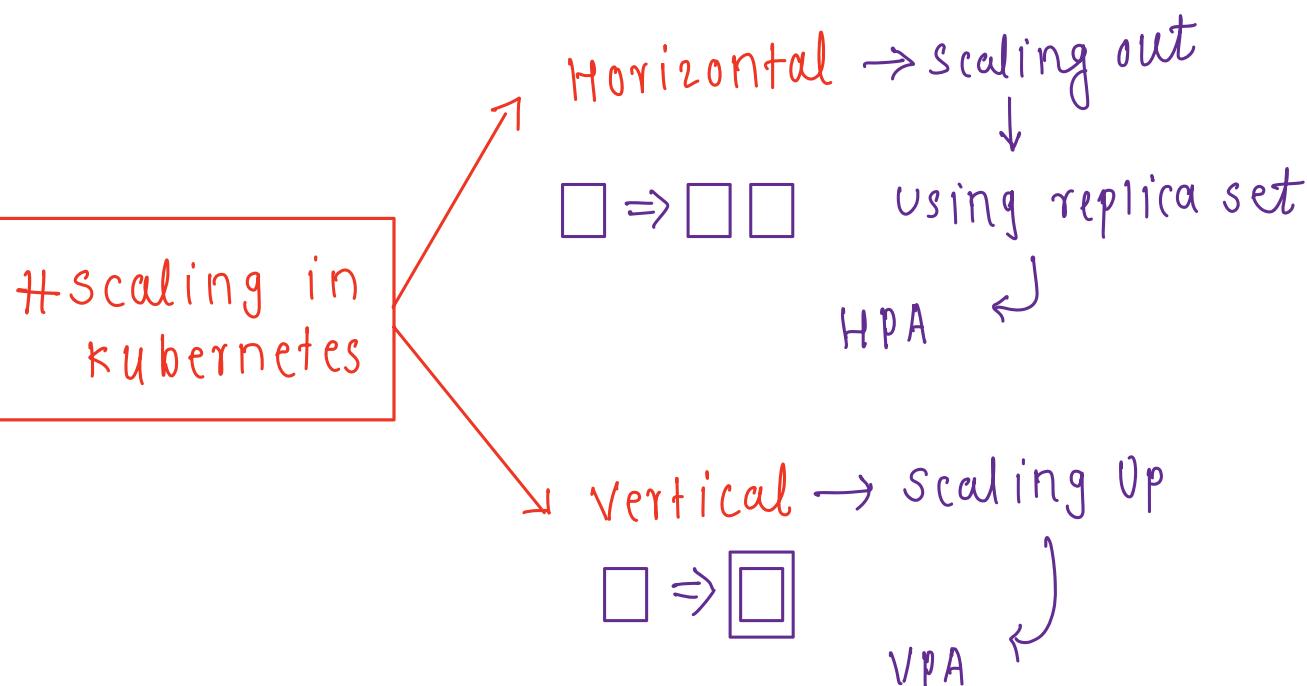
```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  name: deployment-manager
  namespace: deployment
rules:
- apiGroups: ["apps", ""]
  resources: ["deployments"]
  verbs: ["get", "list", "create", "update", "delete"]
---
```

③ Bind role to service Accounts:

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: cicd-deployer-binding
  namespace: deployment
subjects:
- kind: ServiceAccount
  name: cicd-deployer
  namespace: deployment
roleRef:
  kind: Role
  name: deployment-manager
  apiGroup: rbac.authorization.k8s.io
```

④ Attach it to pod:

```
apiVersion: v1
kind: Pod
metadata:
  name: example-pod
spec:
  containers:
  - name: example-container
    image: example/image
  serviceAccountName: example-serviceaccount
```



Kubernetes deployment → High level resource

that manages replica sets.

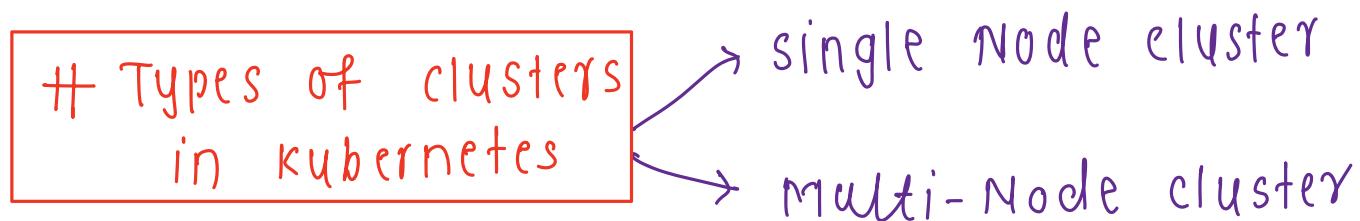
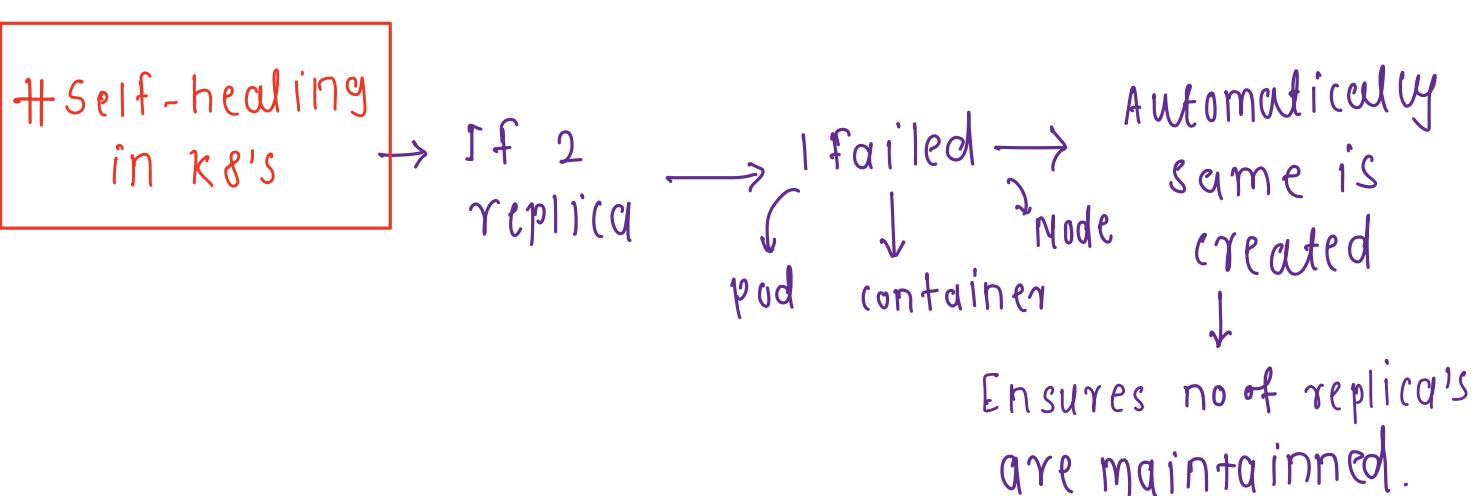


Updates ↘ provides ↓ → Roll back
Rolling update

Kubernetes	Docker Swarm
① complicated setup & strong cluster	① simple setup & weak cluster
② Kubernetes dashboard GUI	② NO GUI
③ Auto scaling	③ No auto-scaling
④ shares storage volume with other containers in same pod	④ shares storage volume with any other containers
⑤ Build in tool for logging & monitoring	⑤ Uses 3rd party tools like ELK
⑥ Load balancing with manual intervention	⑥ automated load balancing

Steps to deploy highly available application using kubernetes

- i) Create kubernetes cluster with 'n' nodes across multiple AZ/regions
- ii) Create K8 deployment → specifies replica
- iii) Create K8 service → provides stable IP & DNS to access applicn
- iv) Configure ingress to load-balance traffic across replica's.



↑↑ kubernetes Namespace

↓
logical clusters inside
kubernetes cluster

system Namespace

- ↑
① default
- ② kube-system
- ③ kube-public
- ④ kube-node-lease

custom Namespace

- ↓
created
- ① dev
- ② QA
- ③ prod

↑↑ storage management in kubernetes.

① volume → way to share data between
containers in pod via filesystem

② configmap → inject config data into pods
(readonly)

③ emptyDir → created when pod is assigned to
node & deleted when pod is deleted. (Temp)

→ created on file system of a node.

→ This directory is shared across all containers
in pod

iii hostpath → file is mounted in pod from node file system.

iv persistentVolumeClaim → request for storage [pvc] made by pod

v secret → to pass sensitive information to pod (readonly)

② persistent volume → provided by administrator at cluster level

③ storage classes → for dynamic provisioning of storage

→ possible to add retain policy

create command → create k8 resource → based on Yaml/json → Error! if resrc already exists

apply command → checks if resrc already exists
create new resrc ✅ update → based on yaml/json

service discovery → how one service found other service internally using dns. [as pod ip changes]

kube-state-metrics → simple service that listens k8 API server → generate metrics about states of the objects.

Service mesh → communication between micro-services
eg. Istio & Linkerd

pod priority and preemption

Assigns priority to pod

Removing low priority pod to give space to high priority pod

Pod preset → inject environment variables

secrets/ volume mounts at pod creation time

Pod scheduler → decides where pod should run → give more control over pods scheduling

Cluster Federation → kubefed → mgmt of a multiple k8's cluster as a single entity.

Resource Quotas → limits resource utilisation

↳ prevents consumption of all cluster resources by any one application

Labels → key value pairs attached to kubernetes resource eg. env=prod , app=frontend

#Selectors → Used to select k8s object based on labels

#Annotations → stores metadata about objects

pod disruption budget [PDB] → k8s object ensures your application remains highly available during voluntary disruptions by limiting how many pods can be down at same time.

- minAvailable → min pods that must stay up
- maxUnavailable → max pods that can be disrupted at once.

#Minikube → runs single-node kubernetes cluster inside VM.

#Skaffold → command line tool → Automate
→ Build → push → deploy in kubernetes

#kompose → migration tool to convert docker compose file into kubernetes manifest, i.e. from local deployments to kubernetes deployment

#kubeadm → Tool for quickly setting up kubernetes cluster using best practices. like installer.

#kernel tuning → adjusting linux parameters on nodes / pods to optimize system performance.

The diagram shows three categories: Network, Memory, and Security. Arrows point from the text 'adjusting linux parameters on nodes / pods to optimize system performance.' to each of these three categories.

#importance of container image security

- i> Avoid vulnerable libraries / Base images
- ii> Eliminate hidden malware
- iii> Secure Embedded secrets (tokens/ SSH keys)
- iv> For image integrity and authenticity
- v> Reduce attack surface for images

#configuration of local storage in k8's / lpv's

① create pv
& apply → ② Deploy pvc → ③ Bind pvc's to pods

CRI [Container Runtime Interface]



Interface required for the communication between kubernetes & container runtimes.

CNI [Container Network interface]



Interface required for the networking in kubernetes

Assign ip addresses to pod

mgmt of network connectivity

mgmt of network policies & Routing

pod Modes Ext network

CSI [Container storage Interface]



Interface used for the mgmt of persistent storage in k8's

AWS EBS
GCP persistent disk

provisioning
Attachment
Mount

Kubernetes Commands & Usages:

1- Basic Cluster Information:

- o `kubectl cluster-info`: Display information about the cluster.
- o `kubectl version`: Display version info.

2- Working with Nodes and Cluster:

- o `kubectl get nodes`: List nodes in a cluster.
- o `kubectl describe node <node-name>`: Show details of a specific node.

3- Working with Pods:

- o `kubectl get pods`: List all pods in all namespaces.
- o `kubectl run <name> --image=<image>`: Deploy a new Pod with a given image.
- o `kubectl describe pod <pod-name>`: Describe a specific pod.
- o `kubectl logs <pod-name>`: Fetch the logs from a pod.
- o `kubectl delete pod <pod-name>`: Delete a specific pod.

4- Working with Deployments:

- o `kubectl create deployment <name> --image=<image>`: Create a new deployment.
- o `kubectl get deployments`: List all deployments.
- o `kubectl describe deployment <deployment-name>`: Describe a specific deployment.
- o `kubectl scale deployment <deployment-name> --replicas=<num>`: Scale up/down a deployment.

5- Working with Services:

- o `kubectl expose deployment <name> --type=LoadBalancer --port=8080`: Expose a deployment as a service.
- o `kubectl get services`: List all services.
- o `kubectl describe service <service-name>`: Describe a specific service.

6- Config and Storage:

- o `kubectl get configmaps`: List all config maps.
- o `kubectl create configmap <name> --from-file=<path>`: Create a config map from a file.
- o `kubectl get secrets`: List all secrets.
- o `kubectl create secret`: Create a secret.
- o `kubectl get pv`: List all persistent volumes.
- o `kubectl get pvc`: List all persistent volume claims.

7- Namespaces and Context:

- o `kubectl get namespaces`: List all namespaces.
- o `kubectl config get-contexts`: Show all contexts.
- o `kubectl config use-context <context-name>`: Switch to a different context.

8- Others:

- o `kubectl apply -f <filename>`: Apply a configuration from a file.
- o `kubectl delete -f <filename>`: Delete resources defined in a file.
- o `kubectl exec -it <pod-name> -- /bin/sh`: Execute a command inside a running pod.
- o `kubectl port-forward <pod-name> <local-port>:<pod-port>`: Forward a port from a running pod to a local port.

9- Advanced:

- o `kubectl get all`: List all resources.
- o `kubectl rollout status deployment/<deployment-name>`: View the rollout status of a deployment.
- o `kubectl rollout history deployment/<deployment-name>`: View the history of a deployment.
- o `kubectl rollout undo deployment/<deployment-name>`: Rollback to a previous version of a deployment.

10- Monitoring & Logging:

- `kubectl top nodes`: Display resource (CPU/Memory/Storage) usage of nodes.
- `kubectl top pods`: Display resource (CPU/Memory/Storage) usage of pods.

11- Autoscaling:

- `kubectl autoscale deployment <deployment-name> --min=<min-pods> --max=<max-pods> --cpu-percent=<cpu-util-percentage>`: Auto scale a deployment based on CPU utilization.

12- Working with Helm:

- `helm list`: List releases.
- `helm install <chart>`: Install a helm chart.
- `helm uninstall <release-name>`: Uninstall a helm release.

13- Using Network Policies:

- `kubectl get networkpolicies`: List all network policies.
- `kubectl describe networkpolicy <policy-name>`: Describe a specific network policy.

14- Using CronJobs and Jobs:

- `kubectl get cronjobs`: List all cronjobs.
- `kubectl get jobs`: List all jobs.
- `kubectl logs job/<job-name>`: Fetch logs from a job.

15- Working with RBAC:

- `kubectl get roles`: List all roles in the current namespace.
- `kubectl get clusterroles`: List all cluster roles.

16- Using kubectl Plugins:

- `kubectl krew search`: Search plugins available for kubectl.
- `kubectl krew install <plugin-name>`: Install a kubectl plugin.

